

Chapter 11

Multithreaded Programming



Key Skills & Concepts

- Understand multithreading fundamentals
- Know the **Thread** class and the **Runnable** interface
- Create a thread
- Create multiple threads
- Determine when a thread ends
- Use thread priorities

- Understand thread synchronization
 - Use synchronized methods
 - Use synchronized blocks
 - Communicate between threads
 - Suspend, resume, and stop threads
-

Although Java contains many innovative features, one of its most exciting is its built-in support for *multithreaded programming*. A multithreaded program contains two or more parts that can run concurrently. Each part of such a program is called a *thread*, and each thread defines a separate path of execution. Thus, multithreading is a specialized form of multitasking.

Multithreading Fundamentals

There are two distinct types of multitasking: process-based and thread-based. It is important to understand the difference between the two. A process is, in essence, a program that is executing. Thus, *process-based* multitasking is the feature that allows your computer to run two or more programs concurrently. For example, it is process-based multitasking that allows you to run the Java compiler at the same time you are using a text editor or browsing the Internet. In process-based multitasking, a program is the smallest unit of code that can be dispatched by the scheduler.

In a *thread-based* multitasking environment, the thread is the smallest unit of dispatchable code. This means that a single program can perform two or more tasks at once. For instance, a text editor can be formatting text at the same time that it is printing, as long as these two actions are being performed by two separate threads. Although Java programs make use of process-based multitasking environments, process-based multitasking is not under the control of Java. Multithreaded multitasking is.

A principal advantage of multithreading is that it enables you to write very efficient programs because it lets you utilize the idle time that is present in most programs. As you probably know, most I/O devices, whether they be network ports, disk drives, or the keyboard, are much slower than the CPU. Thus, a program will often spend a majority of its execution time waiting to send or receive information to or from a device. By using multithreading, your program can execute another task during this idle time. For example, while one part of your program is sending a file over the Internet, another part can be reading keyboard input, and still another can be buffering the next block of data to send.

As you probably know, over the past few years, multiprocessor and multicore systems have become commonplace. Of course, single-processor systems are still in widespread use. It is important to understand that Java's multithreading features work in both types of systems. In a single-core system, concurrently executing threads share the CPU, with each thread receiving a slice of CPU time. Therefore, in a single-core system, two or more threads

do not actually run at the same time, but idle CPU time is utilized. However, in multiprocessor/multicore systems, it is possible for two or more threads to actually execute simultaneously. In many cases, this can further improve program efficiency and increase the speed of certain operations.

A thread can be in one of several states. It can be *running*. It can be *ready to run* as soon as it gets CPU time. A running thread can be *suspended*, which is a temporary halt to its execution. It can later be *resumed*. A thread can be *blocked* when waiting for a resource. A thread can be *terminated*, in which case its execution ends and cannot be resumed.

Along with thread-based multitasking comes the need for a special type of feature called *synchronization*, which allows the execution of threads to be coordinated in certain well-defined ways. Java has a complete subsystem devoted to synchronization, and its key features are also described here.

If you have programmed for operating systems such as Windows, then you are already familiar with multithreaded programming. However, the fact that Java manages threads through language elements makes multithreading especially convenient. Many of the details are handled for you.

The Thread Class and Runnable Interface

Java’s multithreading system is built upon the **Thread** class and its companion interface, **Runnable**. Both are packaged in **java.lang**. **Thread** encapsulates a thread of execution. To create a new thread, your program will either extend **Thread** or implement the **Runnable** interface.

The **Thread** class defines several methods that help manage threads. Here are some of the more commonly used ones (we will be looking at these more closely as they are used):

Method	Meaning
final String getName()	Obtains a thread’s name.
final int getPriority()	Obtains a thread’s priority.
final boolean isAlive()	Determines whether a thread is still running.
final void join()	Waits for a thread to terminate.
void run()	Entry point for the thread.
static void sleep(long milliseconds)	Suspends a thread for a specified period of milliseconds.
void start()	Starts a thread by calling its run() method.

All processes have at least one thread of execution, which is usually called the *main thread*, because it is the one that is executed when your program begins. Thus, the main thread is the thread that all of the preceding example programs in the book have been using. From the main thread, you can create other threads.

Creating a Thread

You create a thread by instantiating an object of type **Thread**. The **Thread** class encapsulates an object that is runnable. As mentioned, Java defines two ways in which you can create a runnable object:

- You can implement the **Runnable** interface.
- You can extend the **Thread** class.

Most of the examples in this chapter will use the approach that implements **Runnable**. However, Try This 11-1 shows how to implement a thread by extending **Thread**. Remember: Both approaches still use the **Thread** class to instantiate, access, and control the thread. The only difference is how a thread-enabled class is created.

The **Runnable** interface abstracts a unit of executable code. You can construct a thread on any object that implements the **Runnable** interface. **Runnable** defines only one method called **run()**, which is declared like this:

```
public void run( )
```

Inside **run()**, you will define the code that constitutes the new thread. It is important to understand that **run()** can call other methods, use other classes, and declare variables just like the main thread. The only difference is that **run()** establishes the entry point for another, concurrent thread of execution within your program. This thread will end when **run()** returns.

After you have created a class that implements **Runnable**, you will instantiate an object of type **Thread** on an object of that class. **Thread** defines several constructors. The one that we will use first is shown here:

```
Thread(Runnable threadOb)
```

In this constructor, *threadOb* is an instance of a class that implements the **Runnable** interface. This defines where execution of the thread will begin.

Once created, the new thread will not start running until you call its **start()** method, which is declared within **Thread**. In essence, **start()** executes a call to **run()**. The **start()** method is shown here:

```
void start( )
```

Here is an example that creates a new thread and starts it running:

```
// Create a thread by implementing Runnable.
```

```
class MyThread implements Runnable {  
    String thrdName;
```

Objects of **MyThread** can be run in their own threads because **MyThread** implements **Runnable**.



```

MyThread(String name) {
    thrdName = name;
}

// Entry point of thread.
public void run() { ←———— Threads start executing here.
    System.out.println(thrdName + " starting.");
    try {
        for(int count=0; count < 10; count++) {
            Thread.sleep(400);
            System.out.println("In " + thrdName +
                               ", count is " + count);
        }
    }
    catch(InterruptedException exc) {
        System.out.println(thrdName + " interrupted.");
    }
    System.out.println(thrdName + " terminating.");
}
}

class UseThreads {
    public static void main(String[] args) {
        System.out.println("Main thread starting.");

        // First, construct a MyThread object.
        MyThread mt = new MyThread("Child #1"); ←———— Create a runnable object.

        // Next, construct a thread from that object.
        Thread newThrd = new Thread(mt); ←———— Construct a thread on that object.

        // Finally, start execution of the thread.
        newThrd.start(); ←———— Start running the thread.

        for(int i=0; i<50; i++) {
            System.out.print(".");
            try {
                Thread.sleep(100);
            }
            catch(InterruptedException exc) {
                System.out.println("Main thread interrupted.");
            }
        }

        System.out.println("Main thread ending.");
    }
}

```


Let's look closely at this program. First, **MyThread** implements **Runnable**. This means that an object of type **MyThread** is suitable for use as a thread and can be passed to the **Thread** constructor.

Inside **run()**, a loop is established that counts from 0 to 9. Notice the call to **sleep()**. The **sleep()** method causes the thread from which it is called to suspend execution for the specified period of milliseconds. Its general form is shown here:

```
static void sleep(long milliseconds) throws InterruptedException
```

The number of milliseconds to suspend is specified in *milliseconds*. This method can throw an **InterruptedException**. Thus, calls to it must be wrapped in a **try** block. The **sleep()** method also has a second form, which allows you to specify the period in terms of milliseconds and nanoseconds if you need that level of precision. In **run()**, **sleep()** pauses the thread for 400 milliseconds each time through the loop. This lets the thread run slow enough for you to watch it execute.

Inside **main()**, a new **Thread** object is created by the following sequence of statements:

```
// First, construct a MyThread object.
MyThread mt = new MyThread("Child #1");

// Next, construct a thread from that object.
Thread newThrd = new Thread(mt);

// Finally, start execution of the thread.
newThrd.start();
```

As the comments suggest, first an object of **MyThread** is created. This object is then used to construct a **Thread** object. This is possible because **MyThread** implements **Runnable**. Finally, execution of the new thread is started by calling **start()**. This causes the child thread's **run()** method to begin. After calling **start()**, execution returns to **main()**, and it enters **main()**'s **for** loop. Notice that this loop iterates 50 times, pausing 100 milliseconds each time through the loop. Both threads continue running, sharing the CPU in single-CPU systems, until their loops finish. The output produced by this program is as follows. Because of differences between computing environments, the precise output that you see may differ slightly from that shown here:

```
Main thread starting.  
.Child #1 starting.  
...In Child #1, count is 0  
....In Child #1, count is 1  
....In Child #1, count is 2  
...In Child #1, count is 3  
....In Child #1, count is 4  
....In Child #1, count is 5  
....In Child #1, count is 6  
...In Child #1, count is 7  
....In Child #1, count is 8  
....In Child #1, count is 9  
Child #1 terminating.  
.....Main thread ending.
```

There is another point of interest to notice in this first threading example. To illustrate the fact that the main thread and **mt** execute concurrently, it is necessary to keep **main()** from terminating until **mt** is finished. Here, this is done through the timing differences between the two threads. Because the calls to **sleep()** inside **main()**'s **for** loop cause a total delay of 5 seconds (50 iterations times 100 milliseconds), but the total delay within **run()**'s loop is only 4 seconds (10 iterations times 400 milliseconds), **run()** will finish approximately 1 second before **main()**. As a result, both the main thread and **mt** will execute concurrently until **mt** ends. Then, about 1 second later **main()** ends.

Although this use of timing differences to ensure that **main()** finishes last is sufficient for this simple example, it is not something that you would normally use in practice. Java provides much better ways of waiting for a thread to end. It is, however, sufficient for the next few programs. Later in this chapter, you will see a better way for one thread to wait until another completes.

Ask the Expert

Q: You state that in a multithreaded program, one will often want the main thread to finish last. Can you explain?

A: The main thread is a convenient place to perform the orderly shutdown of your program, such as the closing of files. It also provides a well-defined exit point for your program. Therefore, it often makes sense for it to finish last. Fortunately, as you will soon see, it is trivially easy for the main thread to wait until the child threads have completed

One other point: In a multithreaded program, you often will want the main thread to be the

last thread to finish running. As a general rule, a program continues to run until all of its threads have ended. Thus, having the main thread finish last is not a requirement. It is, however, often a good practice to follow—especially when you are first learning about threads.

One Improvement and Two Simple Variations

The preceding program demonstrates the fundamentals of creating a **Thread** based on a **Runnable** and then starting the thread. The approach shown in that program is perfectly valid and is often exactly what you will want. However, two simple variations can make **MyThread** more flexible and easier to use in some cases. Furthermore, you may find that these variations are helpful when you create your own **Runnable** classes. It is also possible to make one significant improvement to **MyThread** that takes advantage of another feature of the **Thread** class. Let's begin with the improvement.

In the preceding program, notice that an instance variable called **thrName** is defined by **MyThread** and is used to hold the name of the thread. However, there is no need for **MyThread** to store the name of the thread since it is possible to give a name to a thread when it is created. To do so, use this version of **Thread**'s constructor:

```
Thread(Runnable threadOb, String name)
```

Here, *name* becomes the name of the thread. You can obtain the name of the thread by calling **getName()** defined by **Thread**. Its general form is shown here:

```
final String getName( )
```

Giving a thread a name when it is created provides two advantages. First, there is no need for you to use a separate variable to hold the name because **Thread** already provides this capability. Second, the name of the thread will be available to any code that holds a reference to the thread. One other point: although not needed by this example, you can set the name of a thread after it is created by using **setName()**, which is shown here:

```
final void setName(String threadName)
```

Here, *threadName* specifies the new name of the thread.

As mentioned, there are two variations that can, depending on the situation, make **MyThread** more convenient to use. First, it is possible for the **MyThread** constructor to create a **Thread** object for the thread, storing a reference to that thread in an instance variable. With this approach, the thread is ready to start as soon as the **MyThread** constructor returns. You simply call **start()** on the **Thread** instance encapsulated by **MyThread**.

The second variation offers a way to have a thread begin execution as soon as it is created. This approach is useful in cases in which there is no need to separate thread creation from thread execution. One way to accomplish this for **MyThread** is to provide a static *factory*

method that:

1. creates a new **MyThread** instance,
2. calls **start()** on the thread associated with that instance,
3. and then returns a reference to the newly created **MyThread** object.

With this approach, it becomes possible to create and start a thread through a single method call. This can streamline the use of **MyThread**, especially in cases in which several threads must be created and started.

The following version of the preceding program incorporates the changes just described:

```
// MyThread variations. This version of MyThread
// creates a Thread when its constructor is called and
// stores it in an instance variable called thrd.
// It also sets the name of the thread and provides
// a factory method to create and start a thread.

class MyThread implements Runnable {
    Thread thrd; ←———— A reference to the thread is stored in thrd.

    // Construct a new thread using this Runnable and give
    // it a name.
```

```

MyThread(String name) {
    thrd = new Thread(this, name); ← The thread is named when it is created.
}

// A factory method that creates and starts a thread.
public static MyThread createAndStart(String name) {
    MyThread myThrd = new MyThread(name);

    myThrd.thrd.start(); // start the thread ← Begin executing the thread.
    return myThrd;
}

// Entry point of thread.
public void run() {
    System.out.println(thrd.getName() + " starting.");
    try {
        for(int count=0; count<10; count++) {
            Thread.sleep(400);
            System.out.println("In " + thrd.getName() +
                               ", count is " + count);
        }
    }
    catch(InterruptedException exc) {
        System.out.println(thrd.getName() + " interrupted.");
    }
    System.out.println(thrd.getName() + " terminating.");
}
}

class ThreadVariations {
    public static void main(String[] args) {
        System.out.println("Main thread starting.");

        // Create and start a thread.
        MyThread mt = MyThread.createAndStart("Child #1");

        for(int i=0; i < 50; i++) {
            System.out.print(".");
            try {
                Thread.sleep(100);
            }
            catch(InterruptedException exc) {
                System.out.println("Main thread interrupted.");
            }
        }

        System.out.println("Main thread ending.");
    }
}

```

↑ Now the thread starts when it is created.

This version produces the same output as before. However, notice that now **MyThread** no longer contains the name of the thread. Instead, it provides an instance variable called **thrd** that holds a reference to the **Thread** object created by **MyThread**'s constructor, shown here:

```
MyThread(String name) {  
    thrd = new Thread(this, name);  
}
```

Thus, after **MyThread**'s constructor executes, **thrd** will contain a reference to the newly created thread. To start the thread, you will simply call **start()** on **thrd**.

Next, pay special attention to the **createAndStart()** factory method, shown here:

```
// A factory method that creates and starts a thread.  
public static MyThread createAndStart(String name) {  
    MyThread myThrd = new MyThread(name);  
  
    myThrd.thrd.start(); // start the thread  
    return myThrd;  
}
```

When this method is called, it creates a new instance of **MyThread** called **myThrd**. It then calls **start()** on **myThrd**'s copy of **thrd**. Finally, it returns a reference to the newly created **MyThread** instance. Thus, once the call to **createAndStart()** returns, the thread will already have been started. Therefore, in **main()**, this line creates and begins the execution of a thread in a single call:

```
MyThread mt = MyThread.createAndStart("Child #1");
```

Because of the convenience that **createAndStart()** offers, it will be used by several of the examples in this chapter. Furthermore, you may find it helpful to adapt such a method for use in thread-based applications of your own. Of course, in cases in which you want a thread's execution to be separate from its creation, you can simply create a **MyThread** object and then call **start()** later.

Ask the Expert

Q: Earlier, you used the term *factory method* and showed one example in the method called **createAndStart()**. Can you give me a more general definition?

A: Yes. In general, a factory method is a method that returns an object of a class. Typically, factory methods are **static** methods of a class. Factory methods are useful in a variety of situations. Here are some examples. As you just saw in the case of **createAndStart()**, a factory method enables an object to be constructed

and then set to some specified state prior to being returned to the caller. Another type of factory method is used to provide an easy-to-remember name that indicates the variety of object that is being constructed. For example, assuming a class called **Line**, you might have factory methods that create lines of specific colors, such as **createRedLine()** or **createBlueLine()**. Instead of having to remember a potentially complex call to a constructor, you can simply use the factory method whose name indicates the type of line you want. In some cases it is also possible for a factory method to reuse an object, rather than constructing a new one. As you will see as you advance in your study of Java, factory methods are common in the Java API library.

Try This 11-1 Extending Thread

ExtendThread.java

Implementing **Runnable** is one way to create a class that can instantiate thread objects. Extending **Thread** is the other. In this project, you will see how to extend **Thread** by creating a program functionally similar to the **UseThreads** program shown at the start of this chapter.

When a class extends **Thread**, it must override the **run()** method, which is the entry point for the new thread. It must also call **start()** to begin execution of the new thread. It is possible to override other **Thread** methods, but doing so is not required.

1. Create a file called **ExtendThread.java**. Begin this file with the following lines:

```
/*  
    Try This 11-1  
  
    Extend Thread.  
*/  
class MyThread extends Thread {
```

Notice that **MyThread** now extends **Thread** instead of implementing **Runnable**.

2. Add the following **MyThread** constructor:

```
// Construct a new thread.  
MyThread(String name) {  
    super(name); // name thread  
}
```

Here, **super** is used to call this version of **Thread**'s constructor:

```
Thread(String threadName)
```

Here, *threadName* specifies the name of the thread. As explained previously, **Thread** provides the ability to hold a thread's name. Thus, no instance variable is required by **MyThread** to store the name.

3. Conclude **MyThread** by adding the following **run()** method:

```
// Entry point of thread.
public void run() {

    System.out.println(getName() + " starting.");
    try {
        for(int count=0; count < 10; count++) {
            Thread.sleep(400);
            System.out.println("In " + getName() +
                               ", count is " + count);
        }
    }
    catch(InterruptedException exc) {
        System.out.println(getName() + " interrupted.");
    }

    System.out.println(getName() + " terminating.");
}
```

Notice the calls to **getName()**. Because **ExtendThread** extends **Thread**, it can directly call all of **Thread**'s methods, including the **getName()** method.

4. Next, add the **ExtendThread** class shown here:

```

class ExtendThread {
    public static void main(String[] args) {
        System.out.println("Main thread starting.");

        MyThread mt = new MyThread("Child #1");

        mt.start();

        for(int i=0; i < 50; i++) {
            System.out.print(".");
            try {
                Thread.sleep(100);
            }
            catch(InterruptedException exc) {
                System.out.println("Main thread interrupted.");
            }
        }

        System.out.println("Main thread ending.");
    }
}

```

In **main()**, notice how an instance of **MyThread** is created and then started with these two lines:

```

MyThread mt = new MyThread("Child #1");
mt.start();

```

Because **MyThread** now implements **Thread**, **start()** is called directly on the **MyThread** instance, **mt**.

5. Here is the complete program. Its output is the same as the **UseThreads** example, but in this case, **Thread** is extended rather than **Runnable** being implemented.


```

/*
    Try This 11-1

    Extend Thread.
*/
class MyThread extends Thread {

    // Construct a new thread.
    MyThread(String name) {
        super(name); // name thread
    }

    // Entry point of thread.
    public void run() {
        System.out.println(getName() + " starting.");
        try {
            for(int count=0; count < 10; count++) {
                Thread.sleep(400);
                System.out.println("In " + getName() +
                                   ", count is " + count);
            }
        }
        catch(InterruptedException exc) {
            System.out.println(getName() + " interrupted.");
        }

        System.out.println(getName() + " terminating.");
    }
}

class ExtendThread {
    public static void main(String[] args) {
        System.out.println("Main thread starting.");

        MyThread mt = new MyThread("Child #1");

        mt.start();

        for(int i=0; i < 50; i++) {
            System.out.print(".");
            try {

```

```

        Thread.sleep(100);
    }
    catch (InterruptedException exc) {
        System.out.println("Main thread interrupted.");
    }
}

System.out.println("Main thread ending.");
}
}

```

6. When extending **Thread**, it is also possible to include the ability to create and start a thread in one step by using a **static** factory method, similar to that used by the **ThreadVariations** program shown earlier. To try this, add the following method to **MyThread**:

```

public static MyThread createAndStart(String name) {
    MyThread myThrd = new MyThread(name);
    myThrd.start();
    return myThrd;
}

```

As you can see, this method creates a new **MyThread** instance with the specified name, calls **start()** on that thread, and returns a reference to the thread. To try **createAndStart()**, replace these two lines in **main()**:

```

System.out.println("Main thread starting.");
MyThread mt = new MyThread("Child #1");

```

with this line:

```

MyThread mt = MyThread.createAndStart("Child #1");

```

After making these changes, the program will run the same as before, but you will be creating and starting the thread using a single method call.

Creating Multiple Threads

The preceding examples have created only one child thread. However, your program can spawn as many threads as it needs. For example, the following program creates three child threads:

```
// Create multiple threads.

class MyThread implements Runnable {
    Thread thrd;

    // Construct a new thread.
    MyThread(String name) {
        thrd = new Thread(this, name);
    }
}
```

```

// A factory method that creates and starts a thread.
public static MyThread createAndStart(String name) {
    MyThread myThrd = new MyThread(name);

    myThrd.thrd.start(); // start the thread
    return myThrd;
}

// Entry point of thread.
public void run() {
    System.out.println(thrd.getName() + " starting.");
    try {
        for(int count=0; count < 10; count++) {
            Thread.sleep(400);
            System.out.println("In " + thrd.getName() +
                               ", count is " + count);
        }
    }
    catch(InterruptedException exc) {
        System.out.println(thrd.getName() + " interrupted.");
    }
    System.out.println(thrd.getName() + " terminating.");
}
}

class MoreThreads {
    public static void main(String[] args) {
        System.out.println("Main thread starting.");

        MyThread mt1 = MyThread.createAndStart("Child #1");
        MyThread mt2 = MyThread.createAndStart("Child #2");
        MyThread mt3 = MyThread.createAndStart("Child #3");

        for(int i=0; i < 50; i++) {
            System.out.print(".");
            try {
                Thread.sleep(100);
            }
            catch(InterruptedException exc) {
                System.out.println("Main thread interrupted.");
            }
        }

        System.out.println("Main thread ending.");
    }
}

```

← Create and start
executing three threads.

Ask the Expert

Q: Why does Java have two ways to create child threads (by extending `Thread` or implementing `Runnable`) and which approach is better?

A: The `Thread` class defines several methods that can be overridden by a derived class. Of these methods, the only one that must be overridden is `run()`. This is, of course, the same method required when you implement `Runnable`. Some Java programmers feel that classes should be extended only when they are being expanded or customized in some way. So, if you will not be overriding any of `Thread`'s other methods, it is probably best to simply implement `Runnable`. Also, by implementing `Runnable`, you enable your thread to inherit a class other than `Thread`.

Sample output from this program follows:

```

Main thread starting.
Child #1 starting.
Child #2 starting.
Child #3 starting.
...In Child #3, count is 0
In Child #2, count is 0
In Child #1, count is 0
....In Child #1, count is 1
In Child #2, count is 1
In Child #3, count is 1
....In Child #2, count is 2
In Child #3, count is 2
In Child #1, count is 2
...In Child #1, count is 3
In Child #2, count is 3
In Child #3, count is 3
....In Child #1, count is 4
In Child #3, count is 4
In Child #2, count is 4
....In Child #1, count is 5
In Child #3, count is 5
In Child #2, count is 5
...In Child #3, count is 6
In Child #2, count is 6
In Child #1, count is 6
...In Child #3, count is 7
In Child #1, count is 7
In Child #2, count is 7
....In Child #2, count is 8

In Child #1, count is 8
In Child #3, count is 8
....In Child #1, count is 9
Child #1 terminating.
In Child #2, count is 9
Child #2 terminating.
In Child #3, count is 9
Child #3 terminating.
.....Main thread ending.

```

As you can see, once started, all three child threads share the CPU. Notice that in this run the threads are started in the order in which they are created. However, this may not always be the case. Java is free to schedule the execution of threads in its own way. Of course, because of differences in timing or environment, the precise output from the program may differ, so don't be surprised if you see slightly different results when you try the program.

Determining When a Thread Ends

It is often useful to know when a thread has ended. For example, in the preceding examples, for the sake of illustration it was helpful to keep the main thread alive until the other threads ended. In those examples, this was accomplished by having the main thread sleep longer than the child threads that it spawned. This is, of course, hardly a satisfactory or generalizable solution!

Fortunately, **Thread** provides two means by which you can determine if a thread has ended. First, you can call **isAlive()** on the thread. Its general form is shown here:

```
final boolean isAlive( )
```

The **isAlive()** method returns **true** if the thread upon which it is called is still running. It returns **false** otherwise. To try **isAlive()**, substitute this version of **MoreThreads** for the one shown in the preceding program:

```
// Use isAlive().
class MoreThreads {
    public static void main(String[] args) {
        System.out.println("Main thread starting.");

        MyThread mt1 = MyThread.createAndStart("Child #1");
        MyThread mt2 = MyThread.createAndStart("Child #2");
        MyThread mt3 = MyThread.createAndStart("Child #3");

        do {
            System.out.print(".");
            try {
                Thread.sleep(100);
            }
            catch (InterruptedException exc) {
                System.out.println("Main thread interrupted.");
            }
        } while (mt1.thrd.isAlive() ||
                mt2.thrd.isAlive() || ← This waits until all threads terminate.
                mt3.thrd.isAlive());

        System.out.println("Main thread ending.");
    }
}
```

This version produces output that is similar to the previous version, except that **main()** ends as soon as the other threads finish. The difference is that it uses **isAlive()** to wait for the

child threads to terminate. Another way to wait for a thread to finish is to call **join()**, shown here:

```
final void join( ) throws InterruptedException
```

This method waits until the thread on which it is called terminates. Its name comes from the concept of the calling thread waiting until the specified thread *joins* it. Additional forms of **join()** allow you to specify a maximum amount of time that you want to wait for the specified thread to terminate.

Here is a program that uses **join()** to ensure that the main thread is the last to stop:

```
// Use join().

class MyThread implements Runnable {
    Thread thrd;

    // Construct a new thread.
    MyThread(String name) {
        thrd = new Thread(this, name);
    }

    // A factory method that creates and starts a thread.
    public static MyThread createAndStart(String name) {
        MyThread myThrd = new MyThread(name);

        myThrd.thrd.start(); // start the thread
        return myThrd;
    }

    // Entry point of thread.
    public void run() {
        System.out.println(thrd.getName() + " starting.");
        try {
            for(int count=0; count < 10; count++) {
                Thread.sleep(400);
                System.out.println("In " + thrd.getName() +
                                   ", count is " + count);
            }
        }
    }
}
```

```

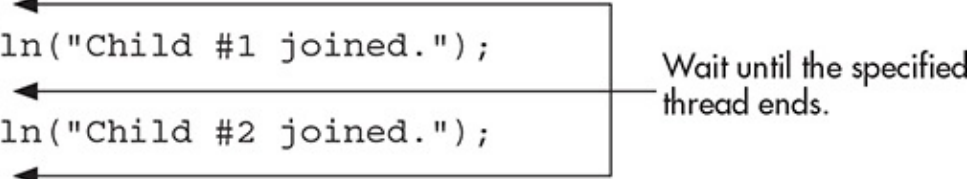
    catch (InterruptedException exc) {
        System.out.println(thrd.getName() + " interrupted.");
    }
    System.out.println(thrd.getName() + " terminating.");
}
}

class JoinThreads {
    public static void main(String[] args) {
        System.out.println("Main thread starting.");

        MyThread mt1 = MyThread.createAndStart("Child #1");
        MyThread mt2 = MyThread.createAndStart("Child #2");
        MyThread mt3 = MyThread.createAndStart("Child #3");

        try {
            mt1.thrd.join(); ←
            System.out.println("Child #1 joined.");
            mt2.thrd.join(); ←
            System.out.println("Child #2 joined.");
            mt3.thrd.join(); ←
            System.out.println("Child #3 joined.");
        }
        catch (InterruptedException exc) {
            System.out.println("Main thread interrupted.");
        }
        System.out.println("Main thread ending.");
    }
}

```



Wait until the specified thread ends.

Sample output from this program is shown here. Remember that when you try the program, your precise output may vary slightly.

```
Main thread starting.  
Child #1 starting.  
Child #2 starting.  
Child #3 starting.  
In Child #2, count is 0  
In Child #1, count is 0  
In Child #3, count is 0  
In Child #2, count is 1  
In Child #3, count is 1  
In Child #1, count is 1  
In Child #2, count is 2  
In Child #1, count is 2  
In Child #3, count is 2  
In Child #2, count is 3  
  
In Child #3, count is 3  
In Child #1, count is 3  
In Child #3, count is 4  
In Child #2, count is 4  
In Child #1, count is 4  
In Child #3, count is 5  
In Child #1, count is 5  
In Child #2, count is 5  
In Child #3, count is 6  
In Child #2, count is 6  
In Child #1, count is 6  
In Child #3, count is 7  
In Child #1, count is 7  
In Child #2, count is 7  
In Child #3, count is 8  
In Child #2, count is 8  
In Child #1, count is 8  
In Child #3, count is 9  
Child #3 terminating.  
In Child #2, count is 9  
Child #2 terminating.  
In Child #1, count is 9  
Child #1 terminating.  
Child #1 joined.  
Child #2 joined.  
Child #3 joined.  
Main thread ending.
```

As you can see, after the calls to **join()** return, the threads have stopped executing.

Thread Priorities

Each thread has associated with it a priority setting. A thread's priority determines, in part, how much CPU time a thread receives relative to the other active threads. In general, over a given period of time, low-priority threads receive little. High-priority threads receive a lot. As you might expect, how much CPU time a thread receives has profound impact on its execution characteristics and its interaction with other threads currently executing in the system.

It is important to understand that factors other than a thread's priority also affect how much CPU time a thread receives. For example, if a high-priority thread is waiting on some resource, perhaps for keyboard input, then it will be blocked, and a lower-priority thread will run. However, when that high-priority thread gains access to the resource, it can preempt the low-priority thread and resume execution. Another factor that affects the scheduling of threads is the way the operating system implements multitasking. (See "Ask the Expert" at the end of this section.) Thus, just because you give one thread a high priority and another a low priority does not necessarily mean that one thread will run faster or more often than the other. It's just that the high-priority thread has greater potential access to the CPU.

When a child thread is started, its priority setting is equal to that of its parent thread. You can change a thread's priority by calling **setPriority()**, which is a member of **Thread**. This is its general form:

```
final void setPriority(int level)
```

Here, *level* specifies the new priority setting for the calling thread. The value of *level* must be within the range **MIN_PRIORITY** and **MAX_PRIORITY**. Currently, these values are 1 and 10, respectively. To return a thread to default priority, specify **NORM_PRIORITY**, which is currently 5. These priorities are defined as **static final** variables within **Thread**.

You can obtain the current priority setting by calling the **getPriority()** method of **Thread**, shown here:

```
final int getPriority( )
```

The following example demonstrates threads at different priorities. The threads are created as instances of **Priority**. The **run()** method contains a loop that counts the number of iterations. The loop stops when either the count reaches 10,000,000 or the static variable **stop** is **true**. Initially, **stop** is set to **false**, but the first thread to finish counting sets **stop** to **true**. This causes each other thread to terminate with its next time slice. Each time through the loop the string in **currentName** is checked against the name of the executing thread. If they don't match, it means that a task-switch occurred. Each time a task-switch happens, the name of the new thread is displayed, and **currentName** is given the name of the new thread. Displaying each thread switch allows you to watch (in a very imprecise way) when the threads gain access to the CPU. After the threads stop, the number of iterations for each loop is displayed.

```
// Demonstrate thread priorities.

class Priority implements Runnable {
    int count;
    Thread thrd;

    static boolean stop = false;
    static String currentName;

    // Construct a new thread.
    Priority(String name) {
        thrd = new Thread(this, name);
        count = 0;
        currentName = name;
    }
}
```



```

// Entry point of thread.
public void run() {
    System.out.println(thrd.getName() + " starting.");
    do {
        count++;

        if (currentName.compareTo(thrd.getName()) != 0) {
            currentName = thrd.getName();
            System.out.println("In " + currentName);
        }

    } while (stop == false && count < 10000000); ← The first thread to
    stop = true;                                10,000,000 stops
                                                all threads.

    System.out.println("\n" + thrd.getName() +
        " terminating.");
}
}

class PriorityDemo {
    public static void main(String[] args) {
        Priority mt1 = new Priority("High Priority");
        Priority mt2 = new Priority("Low Priority");
        Priority mt3 = new Priority("Normal Priority #1");
        Priority mt4 = new Priority("Normal Priority #2");
        Priority mt5 = new Priority("Normal Priority #3");

        // set the priorities
        mt1.thrd.setPriority(Thread.NORM_PRIORITY+2);
        mt2.thrd.setPriority(Thread.NORM_PRIORITY-2); ← Give mt1 a higher priority
        // Leave mt3, mt4, and mt5 at the default, normal priority level
                                                    than mt2.

        // start the threads
        mt1.thrd.start();
        mt2.thrd.start();
        mt3.thrd.start();
        mt4.thrd.start();
        mt5.thrd.start();

        try {
            mt1.thrd.join();
            mt2.thrd.join();
            mt3.thrd.join();
            mt4.thrd.join();
            mt5.thrd.join();
        }
    }
}

```

```

        catch(InterruptedException exc) {
            System.out.println("Main thread interrupted.");
        }

        System.out.println("\nHigh priority thread counted to " +
                           mt1.count);
        System.out.println("Low priority thread counted to " +
                           mt2.count);
        System.out.println("1st Normal priority thread counted to " +
                           mt3.count);
        System.out.println("2nd Normal priority thread counted to " +
                           mt4.count);
        System.out.println("3rd Normal priority thread counted to " +
                           mt5.count);
    }
}

```

Here are the results of a sample run:

```

High priority thread counted to 10000000
Low priority thread counted to 3477862
1st Normal priority thread counted to 7000045
2nd Normal priority thread counted to 6576054
3rd Normal priority thread counted to 7373846

```

In this run, the high-priority thread got the greatest amount of the CPU time. Of course, the exact output produced by this program will depend upon a number of factors, including the speed of your CPU, the number of CPUs in your system, the operating system you are using, and the number and nature of other tasks running in the system. Thus, for any given run, it is actually possible for the low-priority thread to get the most CPU time if the circumstances are right.

Ask the Expert

Q: Does the operating system's implementation of multitasking affect how much CPU time a thread receives?

A: Aside from a thread's priority setting, the most important factor affecting thread execution is the way the operating system implements multitasking and scheduling. Some operating systems use preemptive multitasking in which each thread receives a time slice, at least occasionally. Other systems use nonpreemptive scheduling in which one thread must yield execution before another thread will execute. In

nonpreemptive systems, it is easy for one thread to dominate, preventing others from running.

Synchronization

When using multiple threads, it is sometimes necessary to coordinate the activities of two or more. The process by which this is achieved is called *synchronization*. The most common reason for synchronization is when two or more threads need access to a shared resource that can be used by only one thread at a time. For example, when one thread is writing to a file, a second thread must be prevented from doing so at the same time. Another reason for synchronization is when one thread is waiting for an event that is caused by another thread. In this case, there must be some means by which the first thread is held in a suspended state until the event has occurred. Then, the waiting thread must resume execution.

Key to synchronization in Java is the concept of the *monitor*, which controls access to an object. A monitor works by implementing the concept of a *lock*. When an object is locked by one thread, no other thread can gain access to the object. When the thread exits, the object is unlocked and is available for use by another thread.

All objects in Java have a monitor. This feature is built into the Java language itself. Thus, all objects can be synchronized. Synchronization is supported by the keyword **synchronized** and a few well-defined methods that all objects have. Since synchronization was designed into Java from the start, it is much easier to use than you might first expect. In fact, for many programs, the synchronization of objects is almost transparent.

There are two ways that you can synchronize your code. Both involve the use of the **synchronized** keyword, and both are examined here.

Using Synchronized Methods

You can synchronize access to a method by modifying it with the **synchronized** keyword. When that method is called, the calling thread enters the object's monitor, which then locks the object. While locked, no other thread can enter the method, or enter any other synchronized method defined by the object's class. When the thread returns from the method, the monitor unlocks the object, allowing it to be used by the next thread. Thus, synchronization is achieved with virtually no programming effort on your part.

The following program demonstrates synchronization by controlling access to a method called **sumArray()**, which sums the elements of an integer array.

```
// Use synchronize to control access.

class SumArray {
    private int sum;

    synchronized int sumArray(int[] nums) { ←———— sumArray() is synchronized.
        sum = 0; // reset sum

        for(int i=0; i<nums.length; i++) {
            sum += nums[i];
            System.out.println("Running total for " +
                               Thread.currentThread().getName() +
                               " is " + sum);
        }
    }
}
```

```

        Thread.sleep(10); // allow task-switch
    }
    catch(InterruptedException exc) {
        System.out.println("Thread interrupted.");
    }
}
return sum;
}
}

class MyThread implements Runnable {
    Thread thrd;
    static SumArray sa = new SumArray();
    int[] a;
    int answer;

    // Construct a new thread.
    MyThread(String name, int[] nums) {
        thrd = new Thread(this, name);
        a = nums;
    }

    // A factory method that creates and starts a thread.
    public static MyThread createAndStart(String name, int[] nums) {
        MyThread myThrd = new MyThread(name, nums);

        myThrd.thrd.start(); // start the thread
        return myThrd;
    }

    // Entry point of thread.
    public void run() {
        int sum;

        System.out.println(thrd.getName() + " starting.");

        answer = sa.sumArray(a);
        System.out.println("Sum for " + thrd.getName() +
            " is " + answer);

        System.out.println(thrd.getName() + " terminating.");
    }
}

class Sync {
    public static void main(String[] args) {
        int[] a = {1, 2, 3, 4, 5};

        MyThread mt1 = MyThread.createAndStart("Child #1", a);
        MyThread mt2 = MyThread.createAndStart("Child #2", a);
    }
}

```



```

    try {
        mt1.thrd.join();
        mt2.thrd.join();
    }
    catch (InterruptedException exc) {
        System.out.println("Main thread interrupted.");
    }
}
}

```

The output from the program is shown here. (The precise output may differ on your computer.)

```

Child #1 starting.
Running total for Child #1 is 1
Child #2 starting.
Running total for Child #1 is 3
Running total for Child #1 is 6
Running total for Child #1 is 10
Running total for Child #1 is 15
Sum for Child #1 is 15
Child #1 terminating.
Running total for Child #2 is 1
Running total for Child #2 is 3
Running total for Child #2 is 6
Running total for Child #2 is 10
Running total for Child #2 is 15
Sum for Child #2 is 15
Child #2 terminating.

```

Let's examine this program in detail. The program creates three classes. The first is **SumArray**. It contains the method **sumArray()**, which sums an integer array. The second class is **MyThread**, which uses a **static** object of type **SumArray** to obtain the sum of an integer array. This object is called **sa** and because it is **static**, there is only one copy of it that is shared by all instances of **MyThread**. Finally, the class **Sync** creates two threads and has each compute the sum of an integer array.

Inside **sumArray()**, **sleep()** is called to purposely allow a task switch to occur, if one can—but it can't. Because **sumArray()** is synchronized, it can be used by only one thread at a time. Thus, when the second child thread begins execution, it does not enter **sumArray()** until after the first child thread is done with it. This ensures that the correct result is produced.

To fully understand the effects of **synchronized**, try removing it from the declaration of **sumArray()**. After doing this, **sumArray()** is no longer synchronized, and any number of

threads may use it concurrently. The problem with this is that the running total is stored in **sum**, which will be changed by each thread that calls **sumArray()** through the **static** object **sa**. Thus, when two threads call **sa.sumArray()** at the same time, incorrect results are produced because **sum** reflects the summation of both threads, mixed together. For example, here is sample output from the program after **synchronized** has been removed from **sumArray()**'s declaration. (The precise output may differ on your computer.)

```
Child #1 starting.  
Running total for Child #1 is 1  
Child #2 starting.  
Running total for Child #2 is 1  
Running total for Child #1 is 3  
Running total for Child #2 is 5  
Running total for Child #2 is 8  
Running total for Child #1 is 11  
Running total for Child #2 is 15  
Running total for Child #1 is 19  
Running total for Child #2 is 24  
Sum for Child #2 is 24  
Child #2 terminating.  
Running total for Child #1 is 29  
Sum for Child #1 is 29  
Child #1 terminating.
```

As the output shows, both child threads are calling **sa.sumArray()** concurrently, and the value of **sum** is corrupted. Before moving on, let's review the key points of a synchronized method:

- A synchronized method is created by preceding its declaration with **synchronized**.
- For any given object, once a synchronized method has been called, the object is locked and no synchronized methods on the same object can be used by another thread of execution.
- Other threads trying to call an in-use synchronized object will enter a wait state until the object is unlocked.
- When a thread leaves the synchronized method, the object is unlocked.

The synchronized Statement

Although creating **synchronized** methods within classes that you create is an easy and effective means of achieving synchronization, it will not work in all cases. For example, you might want to synchronize access to some method that is not modified by **synchronized**. This can occur because you want to use a class that was not created by you but by a third party, and you do not have access to the source code. Thus, it is not possible for you to add

synchronized to the appropriate methods within the class. How can access to an object of this class be synchronized? Fortunately, the solution to this problem is quite easy: You simply put calls to the methods defined by this class inside a **synchronized** block.

This is the general form of a **synchronized** block:

```
synchronized(objref) {  
    // statements to be synchronized  
}
```

Here, *objref* is a reference to the object being synchronized. Once a synchronized block has been entered, no other thread can call a synchronized method on the object referred to by *objref* until the block has been exited.

For example, another way to synchronize calls to **sumArray()** is to call it from within a synchronized block, as shown in this version of the program:

// Use a synchronized block to control access to SumArray.

```
class SumArray {
    private int sum;

    int sumArray(int[] nums) { ← Here, sumArray()
        sum = 0; // reset sum      is not synchronized.

        for(int i=0; i<nums.length; i++) {
            sum += nums[i];
            System.out.println("Running total for " +
                Thread.currentThread().getName() +
                " is " + sum);
            try {
                Thread.sleep(10); // allow task-switch
            }
            catch(InterruptedException exc) {
                System.out.println("Thread interrupted.");
            }
        }
        return sum;
    }
}
```

```
class MyThread implements Runnable {
    Thread thrd;
    static SumArray sa = new SumArray();
    int[] a;
    int answer;

    // Construct a new thread.
    MyThread(String name, int[] nums) {
        thrd = new Thread(this, name);
        a = nums;
    }

    // A factory method that creates and starts a thread.
    public static MyThread createAndStart(String name, int[] nums) {
        MyThread myThrd = new MyThread(name, nums);

        myThrd.thrd.start(); // start the thread
        return myThrd;
    }
}
```

```

// Entry point of thread.
public void run() {
    int sum;

    System.out.println(thrd.getName() + " starting.");

    // synchronize calls to sumArray()
    synchronized(sa) { ←————— Here, calls to sumArray()
        answer = sa.sumArray(a);      on sa are synchronized.
    }
    System.out.println("Sum for " + thrd.getName() +
        " is " + answer);

    System.out.println(thrd.getName() + " terminating.");
}
}

class Sync {
    public static void main(String[] args) {
        int[] a = {1, 2, 3, 4, 5};

        MyThread mt1 = MyThread.createAndStart("Child #1", a);
        MyThread mt2 = MyThread.createAndStart("Child #2", a);

        try {
            mt1.thrd.join();
            mt2.thrd.join();
        } catch (InterruptedException exc) {
            System.out.println("Main thread interrupted.");
        }
    }
}

```

This version produces the same, correct output as the one shown earlier that uses a synchronized method.

Ask the Expert

Q: I have heard of something called the “concurrency utilities.” What are these? Also, what is the Fork/Join Framework?

A: The concurrency utilities, which are packaged in **java.util.concurrent** (and its

subpackages), support concurrent programming. Among several other items, they offer synchronizers, thread pools, execution managers, and locks that expand your control over thread execution. One of the most exciting features of the concurrent API is the Fork/Join Framework.

The Fork/Join Framework supports what is often termed *parallel programming*. This is the name commonly given to the techniques that take advantage of computers that contain two or more processors (including multicore systems) by subdividing a task into subtasks, with each subtask executing on its own processor. As you can imagine, such an approach can lead to significantly higher throughput and performance. The key advantage of the Fork/Join Framework is ease of use; it streamlines the development of multithreaded code that automatically scales to utilize the number of processors in a system. Thus, it facilitates the creation of concurrent solutions to some common programming tasks, such as performing operations on the elements of an array. The concurrency utilities in general, and the Fork/Join Framework specifically, are features that you will want to explore after you have become more experienced with multithreading.

Thread Communication Using `notify()`, `wait()`, and `notifyAll()`

Consider the following situation. A thread called `T` is executing inside a synchronized method and needs access to a resource called `R` that is temporarily unavailable. What should `T` do? If `T` enters some form of polling loop that waits for `R`, `T` ties up the object, preventing other threads' access to it. This is a less than optimal solution because it partially defeats the advantages of programming for a multithreaded environment. A better solution is to have `T` temporarily relinquish control of the object, allowing another thread to run. When `R` becomes available, `T` can be notified and resume execution. Such an approach relies upon some form of interthread communication in which one thread can notify another that it is blocked and be notified that it can resume execution. Java supports interthread communication with the **`wait()`**, **`notify()`**, and **`notifyAll()`** methods.

The **`wait()`**, **`notify()`**, and **`notifyAll()`** methods are part of all objects because they are implemented by the **`Object`** class. These methods should be called only from within a **synchronized** context. Here is how they are used. When a thread is temporarily blocked from running, it calls **`wait()`**. This causes the thread to go to sleep and the monitor for that object to be released, allowing another thread to use the object. At a later point, the sleeping thread is awakened when some other thread enters the same monitor and calls **`notify()`**, or **`notifyAll()`**.

Following are the various forms of **`wait()`** defined by **`Object`**:

```
final void wait() throws InterruptedException
```

```
final void wait(long millis) throws InterruptedException
```

`final void wait(long millis, int nanos)` throws `InterruptedException`

The first form waits until notified. The second form waits until notified or until the specified period of milliseconds has expired. The third form allows you to specify the wait period in terms of nanoseconds.

Here are the general forms for **notify()** and **notifyAll()**:

```
final void notify( )
```

```
final void notifyAll( )
```

A call to **notify()** resumes one waiting thread. A call to **notifyAll()** notifies all threads, with the scheduler determining which thread gains access to the object.

Before looking at an example that uses **wait()**, an important point needs to be made. Although **wait()** normally waits until **notify()** or **notifyAll()** is called, there is a possibility that in very rare cases the waiting thread could be awakened due to a *spurious wakeup*. The conditions that lead to a spurious wakeup are complex and beyond the scope of this book. However, the Java API documentation recommends that because of the remote possibility of a spurious wakeup, calls to **wait()** should take place within a loop that checks the condition on which the thread is waiting. The following example shows this technique.

An Example That Uses **wait()** and **notify()**

To understand the need for and the application of **wait()** and **notify()**, we will create a program that simulates the ticking of a clock by displaying the words Tick and Tock on the screen. To accomplish this, we will create a class called **TickTock** that contains two methods: **tick()** and **tock()**. The **tick()** method displays the word "Tick", and **tock()** displays "Tock". To run the clock, two threads are created, one that calls **tick()** and one that calls **tock()**. The goal is to make the two threads execute in a way that the output from the program displays a consistent "Tick Tock"—that is, a repeated pattern of one tick followed by one tock.


```
// Use wait() and notify() to create a ticking clock.

class TickTock {

    String state; // contains the state of the clock

    synchronized void tick(boolean running) {
        if(!running) { // stop the clock
            state = "ticked";
            notify(); // notify any waiting threads
            return;
        }

        System.out.print("Tick ");

        state = "ticked"; // set the current state to ticked

        notify(); // let tock() run ←—— tick() notifies tock().
        try {
            while(!state.equals("tocked"))
                wait(); // wait for tock() to complete ←—— tick() waits for tock().
        }
    }
}
```

```

        catch(InterruptedException exc) {
            System.out.println("Thread interrupted.");
        }
    }

    synchronized void tock(boolean running) {
        if(!running) { // stop the clock
            state = "tocked";
            notify(); // notify any waiting threads
            return;
        }

        System.out.println("Tock");

        state = "tocked"; // set the current state to tocked

        notify(); // let tick() run ← tock() notifies tick().
        try {
            while(!state.equals("ticked"))
                wait(); // wait for tick to complete ← tock() waits for tick().
        }
        catch(InterruptedException exc) {
            System.out.println("Thread interrupted.");
        }
    }
}

class MyThread implements Runnable {
    Thread thrd;
    TickTock ttOb;

    // Construct a new thread.
    MyThread(String name, TickTock tt) {
        thrd = new Thread(this, name);
        ttOb = tt;
    }

    // A factory method that creates and starts a thread.
    public static MyThread createAndStart(String name, TickTock tt) {
        MyThread myThrd = new MyThread(name, tt);

        myThrd.thrd.start(); // start the thread
        return myThrd;
    }

    // Entry point of thread.
    public void run() {

```

```

        if(thrd.getName().compareTo("Tick") == 0) {
            for(int i=0; i<5; i++) ttOb.tick(true);
            ttOb.tick(false);
        }
        else {
            for(int i=0; i<5; i++) ttOb.tock(true);
            ttOb.tock(false);
        }
    }
}

class ThreadCom {
    public static void main(String[] args) {
        TickTock tt = new TickTock();
        MyThread mt1 = MyThread.createAndStart("Tick", tt);
        MyThread mt2 = MyThread.createAndStart("Tock", tt);

        try {
            mt1.thrd.join();
            mt2.thrd.join();
        } catch (InterruptedException exc) {
            System.out.println("Main thread interrupted.");
        }
    }
}

```

Here is the output produced by the program:

```

Tick Tock
Tick Tock
Tick Tock
Tick Tock
Tick Tock

```

Let's take a close look at this program. The heart of the clock is the **TickTock** class. It contains two methods, **tick()** and **tock()**, which communicate with each other to ensure that a Tick is always followed by a Tock, which is always followed by a Tick, and so on. Notice the **state** field. When the clock is running, **state** will hold either the string "ticked" or "tocked", which indicates the current state of the clock. In **main()**, a **TickTock** object called **tt** is created, and this object is used to start two threads of execution.

The threads are based on objects of type **MyThread**. Both the **MyThread** constructor and the **createAndStart()** method are passed two arguments. The first becomes the name of the thread. This will be either "Tick" or "Tock". The second is a reference to the **TickTock** object, which is **tt** in this case. Inside the **run()** method of **MyThread**, if the name of the

thread is "Tick", then calls to **tick()** are made. If the name of the thread is "Tock", then the **tock()** method is called. Five calls that pass **true** as an argument are made to each method. The clock runs as long as **true** is passed. A final call that passes **false** to each method stops the clock.

The most important part of the program is found in the **tick()** and **tock()** methods of **TickTock**. We will begin with the **tick()** method, which, for convenience, is shown here:

```
synchronized void tick(boolean running) {
    if(!running) { // stop the clock
        state = "ticked";
        notify(); // notify any waiting threads
        return;
    }

    System.out.print("Tick ");

    state = "ticked"; // set the current state to ticked

    notify(); // let tock() run
    try {
        while(!state.equals("tocked"))
            wait(); // wait for tock() to complete
    }
    catch(InterruptedException exc) {
        System.out.println("Thread interrupted.");
    }
}
```

First, notice that **tick()** is modified by **synchronized**. Remember, **wait()** and **notify()** apply only to synchronized methods. The method begins by checking the value of the **running** parameter. This parameter is used to provide a clean shutdown of the clock. If it is **false**, then the clock has been stopped. If this is the case, **state** is set to "ticked" and a call to **notify()** is made to enable any waiting thread to run. We will return to this point in a moment.

Assuming that the clock is running when **tick()** executes, the word "Tick" is displayed, **state** is set to "ticked", and then a call to **notify()** takes place. The call to **notify()** allows a thread waiting on the same object to run. Next, **wait()** is called within a **while** loop. The call to **wait()** causes **tick()** to suspend until another thread calls **notify()**. Therefore, the loop will not iterate until another thread calls **notify()** on the same object. As a result, when **tick()** is called, it displays one "Tick", lets another thread run, and then suspends.

The **while** loop that calls **wait()** checks the value of **state**, waiting for it to equal "tocked", which will be the case only after the **tock()** method executes. As explained, using a **while** loop to check this condition prevents a spurious wakeup from incorrectly restarting the thread. If **state** does not equal "tocked" when **wait()** returns, it means that a spurious wakeup occurred, and **wait()** is simply called again.

The **tock()** method is an exact copy of **tick()** except that it displays "Tock" and sets **state** to "tocked". Thus, when entered, it displays "Tock", calls **notify()**, and then waits. When viewed as a pair, a call to **tick()** can only be followed by a call to **tock()**, which can only be followed by a call to **tick()**, and so on. Therefore, the two methods are mutually synchronized.

The reason for the call to **notify()** when the clock is stopped is to allow a final call to **wait()** to succeed. Remember, both **tick()** and **tock()** execute a call to **wait()** after displaying their message. The problem is that when the clock is stopped, one of the methods will still be waiting. Thus, a final call to **notify()** is required in order for the waiting method to run. As an experiment, try removing this call to **notify()** and watch what happens. As you will see, the program will "hang," and you will need to press **CTRL-C** to exit. The reason for this is that when the final call to **tock()** calls **wait()**, there is no corresponding call to **notify()** that lets **tock()** conclude. Thus, **tock()** just sits there, waiting forever.

Before moving on, if you have any doubt that the calls to **wait()** and **notify()** are actually needed to make the "clock" run right, substitute this version of **TickTock** into the preceding program. It has all calls to **wait()** and **notify()** removed.

```
// No calls to wait() or notify().
class TickTock {

    String state; // contains the state of the clock

    synchronized void tick(boolean running) {
        if(!running) { // stop the clock
            state = "ticked";
            return;
        }

        System.out.print("Tick ");

        state = "ticked"; // set the current state to ticked
    }

    synchronized void tock(boolean running) {
        if(!running) { // stop the clock
            state = "tocked";
            return;
        }

        System.out.println("Tock");

        state = "tocked"; // set the current state to tocked
    }
}
```

After the substitution, the output produced by the program will look like this:

```
Tick Tick Tick Tick Tick Tock  
Tock  
Tock  
Tock  
Tock
```

Clearly, the `tick()` and `tock()` methods are no longer working together!

Ask the Expert

Q: I have heard the term *deadlock* applied to misbehaving multithreaded programs. What is it, and how can I avoid it? Also, what is a *race condition*, and how can I avoid that, too?

A: Deadlock is, as the name implies, a situation in which one thread is waiting for another thread to do something, but that other thread is waiting on the first. Thus, both threads are suspended, waiting on each other, and neither executes. This situation is analogous to two overly polite people, both insisting that the other step through a door first!

Avoiding deadlock seems easy, but it's not. For example, deadlock can occur in roundabout ways. The cause of the deadlock often is not readily understood just by looking at the source code to the program because concurrently executing threads can interact in complex ways at run time. To avoid deadlock, careful programming and thorough testing is required. Remember, if a multithreaded program occasionally "hangs," deadlock is the likely cause.

A race condition occurs when two (or more) threads attempt to access a shared resource at the same time, without proper synchronization. For example, one thread may be writing a new value to a variable while another thread is incrementing the variable's current value. Without synchronization, the new value of the variable will depend upon the order in which the threads execute. (Does the second thread increment the original value or the new value written by the first thread?) In situations like this, the two threads are said to be "racing each other," with the final outcome determined by which thread finishes first. Like deadlock, a race condition can occur in difficult-to-discover ways. The solution is prevention: careful programming that properly synchronizes access to shared resources.

Suspending, Resuming, and Stopping Threads

It is sometimes useful to suspend execution of a thread. For example, a separate thread can be

used to display the time of day. If the user does not desire a clock, then its thread can be suspended. Whatever the case, it is a simple matter to suspend a thread. Once suspended, it is also a simple matter to restart the thread.

The mechanisms to suspend, stop, and resume threads differ between early versions of Java and more modern versions, beginning with Java 2. Prior to Java 2, a program used **suspend()**, **resume()**, and **stop()**, which are methods defined by **Thread**, to pause, restart, and stop the execution of a thread. They have the following forms:

```
final void resume( )
```

```
final void suspend( )
```

```
final void stop( )
```

While these methods seem to be a perfectly reasonable and convenient approach to managing the execution of threads, they must no longer be used. Here's why. The **suspend()** method of the **Thread** class was deprecated by Java 2. This was done because **suspend()** can sometimes cause serious problems that involve deadlock. The **resume()** method is also deprecated. It does not cause problems but cannot be used without the **suspend()** method as its counterpart. The **stop()** method of the **Thread** class was also deprecated by Java 2. This was done because this method too can sometimes cause serious problems.

Since you cannot now use the **suspend()**, **resume()**, or **stop()** methods to control a thread, you might at first be thinking that there is no way to pause, restart, or terminate a thread. But, fortunately, this is not true. Instead, a thread must be designed so that the **run()** method periodically checks to determine if that thread should suspend, resume, or stop its own execution. Typically, this is accomplished by establishing two flag variables: one for suspend and resume, and one for stop. For suspend and resume, as long as the flag is set to "running," the **run()** method must continue to let the thread execute. If this variable is set to "suspend," the thread must pause. For the stop flag, if it is set to "stop," the thread must terminate.

The following example shows one way to implement your own versions of **suspend()**, **resume()**, and **stop()**:

```
// Suspending, resuming, and stopping a thread.
```

```
class MyThread implements Runnable {
    Thread thrd;
    boolean suspended; ← Suspends thread when true.
    boolean stopped; ← Stops thread when true.

    MyThread(String name) {
        thrd = new Thread(this, name);
        suspended = false;
        stopped = false;
    }

    // A factory method that creates and starts a thread.
    public static MyThread createAndStart(String name) {
        MyThread myThrd = new MyThread(name);

        myThrd.thrd.start(); // start the thread
        return myThrd;
    }

    // Entry point of thread.
    public void run() {
        System.out.println(thrd.getName() + " starting.");
        try {
            for(int i = 1; i < 1000; i++) {
                System.out.print(i + " ");
                if((i%10)==0) {
                    System.out.println();
                    Thread.sleep(250);
                }
            }
        }
    }
}
```

```

    // Use synchronized block to check suspended and stopped.
    synchronized(this) {
        while(suspended) {
            wait();
        }
        if(stopped) break;
    }
} catch (InterruptedException exc) {
    System.out.println(thrd.getName() + " interrupted.");
}
System.out.println(thrd.getName() + " exiting.");
}

// Stop the thread.
synchronized void mystop() {
    stopped = true;

    // The following ensures that a suspended thread can be stopped.
    suspended = false;
    notify();
}

// Suspend the thread.
synchronized void mysuspend() {
    suspended = true;
}

// Resume the thread.
synchronized void myresume() {
    suspended = false;
    notify();
}

}

class Suspend {
    public static void main(String[] args) {
        MyThread mt1 = MyThread.createAndStart("My Thread");

        try {
            Thread.sleep(1000); // let ob1 thread start executing

            mt1.mysuspend();
            System.out.println("Suspending thread.");
            Thread.sleep(1000);

            mt1.myresume();
            System.out.println("Resuming thread.");
            Thread.sleep(1000);
        }
    }
}

```

This synchronized block checks **suspended** and **stopped**.

```

        mt1.mysuspend();
        System.out.println("Suspending thread.");
        Thread.sleep(1000);

        mt1.myresume();
        System.out.println("Resuming thread.");
        Thread.sleep(1000);

        mt1.mysuspend();
        System.out.println("Stopping thread.");
        mt1.mystop();

    } catch (InterruptedException e) {
        System.out.println("Main thread Interrupted");
    }

    // wait for thread to finish
    try {
        mt1.thrd.join();
    } catch (InterruptedException e) {
        System.out.println("Main thread Interrupted");
    }

    System.out.println("Main thread exiting.");
}
}

```

Sample output from this program is shown here. (Your output may differ slightly.)

```
My Thread starting.
1 2 3 4 5 6 7 8 9 10
11 12 13 14 15 16 17 18 19 20
21 22 23 24 25 26 27 28 29 30
31 32 33 34 35 36 37 38 39 40
Suspending thread.
Resuming thread.
41 42 43 44 45 46 47 48 49 50
51 52 53 54 55 56 57 58 59 60
61 62 63 64 65 66 67 68 69 70
71 72 73 74 75 76 77 78 79 80
Suspending thread.
Resuming thread.
81 82 83 84 85 86 87 88 89 90
91 92 93 94 95 96 97 98 99 100
101 102 103 104 105 106 107 108 109 110
111 112 113 114 115 116 117 118 119 120
Stopping thread.
My Thread exiting.
Main thread exiting.
```

Ask the Expert

Q: Multithreading seems like a great way to improve the efficiency of my programs. Can you give me any tips on effectively using it?

A: The key to effectively utilizing multithreading is to think concurrently rather than serially. For example, when you have two subsystems within a program that are fully independent of each other, consider making them into individual threads. A word of caution is in order, however. If you create too many threads, you can actually degrade the performance of your program rather than enhance it. Remember, overhead is associated with context switching. If you create too many threads, more CPU time will be spent changing contexts than in executing your program!

Here is how the program works. The thread class **MyThread** defines two Boolean variables, **suspended** and **stopped**, which govern the suspension and termination of a thread. Both are initialized to **false** by the constructor. The **run()** method contains a **synchronized** statement block that checks **suspended**. If that variable is **true**, the **wait()** method is invoked to suspend the execution of the thread. To suspend execution of the thread, call **mysuspend()**, which sets **suspended** to **true**. To resume execution, call **myresume()**, which sets

suspended to **false** and invokes **notify()** to restart the thread.

To stop the thread, call **mystop()**, which sets **stopped** to **true**. In addition, **mystop()** sets **suspended** to **false** and then calls **notify()**. These steps are necessary to stop a suspended thread.

Try This 11-2 Using the Main Thread

UseMain.jav

All Java programs have at least one thread of execution, called the *main thread*, which is given to the program automatically when it begins running. So far, we have been taking the main thread for granted. In this project, you will see that the main thread can be handled just like all other threads.

1. Create a file called **UseMain.java**.
2. To access the main thread, you must obtain a **Thread** object that refers to it. You do this by calling the **currentThread()** method, which is a **static** member of **Thread**. Its general form is shown here:

```
static Thread currentThread( )
```

This method returns a reference to the thread in which it is called. Therefore, if you call **currentThread()** while execution is inside the main thread, you will obtain a reference to the main thread. Once you have this reference, you can control the main thread just like any other thread.

3. Enter the following program into the file. It obtains a reference to the main thread, and then gets and sets the main thread's name and priority.


```

/*
    Try This 11-2

    Controlling the main thread.
*/

class UseMain {
    public static void main(String[] args) {
        Thread thrd;

        // Get the main thread.
        thrd = Thread.currentThread();

        // Display main thread's name.
        System.out.println("Main thread is called: " +
                           thrd.getName());

        // Display main thread's priority.
        System.out.println("Priority: " +
                           thrd.getPriority());

        System.out.println();

        // Set the name and priority.
        System.out.println("Setting name and priority.\n");
        thrd.setName("Thread #1");
        thrd.setPriority(Thread.NORM_PRIORITY+3);

        System.out.println("Main thread is now called: " +
                           thrd.getName());

        System.out.println("Priority is now: " +
                           thrd.getPriority());
    }
}

```

4. The output from the program is shown here:

```
Main thread is called: main  
Priority: 5
```

Setting name and priority.

```
Main thread is now called: Thread #1  
Priority is now: 8
```

You need to be careful about what operations you perform on the main thread. For example, if you add the following code to the end of **main()**, the program will never terminate because it will be waiting for the main thread to end!

```
try {  
    thrd.join();  
} catch (InterruptedException exc) {  
    System.out.println("Interrupted");  
}
```

✓ Chapter 11 Self Test

1. How does Java's multithreading capability enable you to write more efficient programs?
2. Multithreading is supported by the _____ class and the _____ interface.
3. When creating a runnable object, why might you want to extend **Thread** rather than implement **Runnable**?
4. Show how to use **join()** to wait for a thread object called **MyThrd** to end.
5. Show how to set a thread called **MyThrd** to three levels above normal priority.
6. What is the effect of adding the **synchronized** keyword to a method?
7. The **wait()** and **notify()** methods are used to perform _____.
8. Change the **TickTock** class so that it actually keeps time. That is, have each tick take one half second, and each tock take one half second. Thus, each tick-tock will take one second. (Don't worry about the time it takes to switch tasks, etc.)
9. Why can't you use **suspend()**, **resume()**, and **stop()** for new programs?
10. What method defined by **Thread** obtains the name of a thread?
11. What does **isAlive()** return?
12. On your own, try adding synchronization to the **Queue** class developed in previous chapters so that it is safe for multithreaded use.